Maintained benchmark task/tool state transfer

164 canonical tasks × four controls = 656 planned cells; native tools and DB evaluator unchanged.

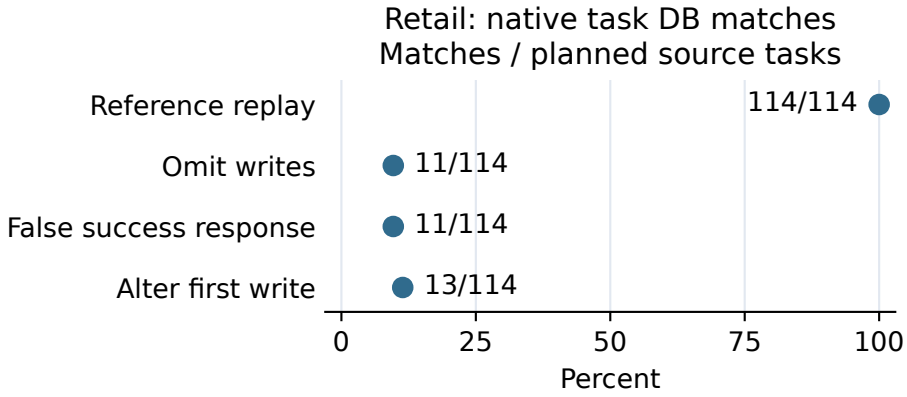

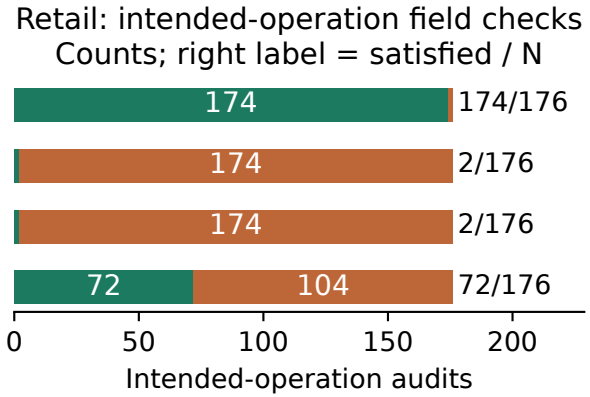

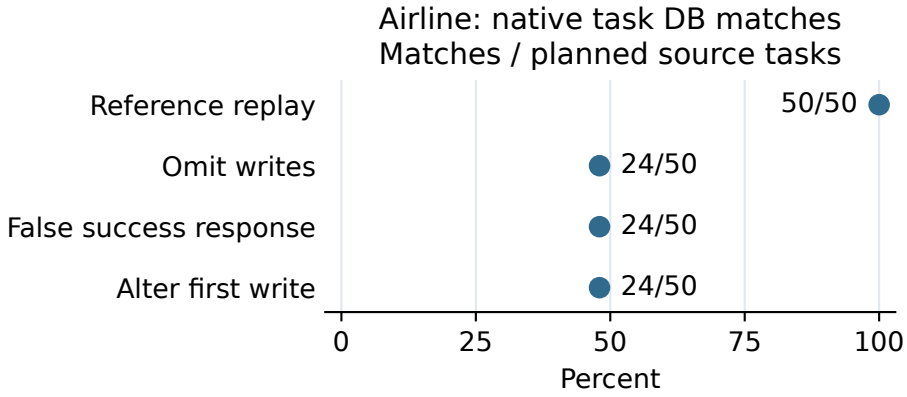

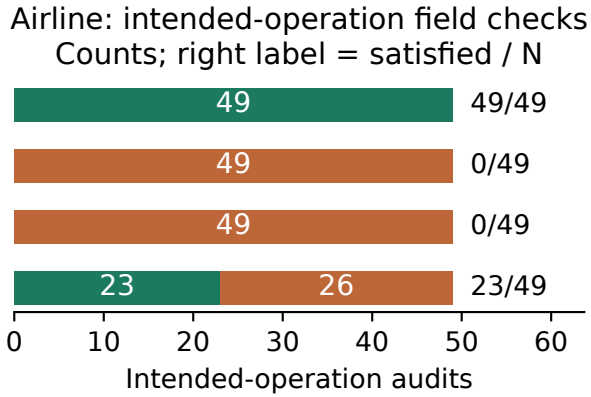

Model-free reference/control replay, not autonomous-agent benchmark performance.
DB targets and immediate field contracts differ; these counts are not predicate accuracy.
Missing readback and origin remain unknown.